Same tools. Different results.

Your company has rolled out the same AI as everyone else. Inside, results vary person by person, and so does what people believe AI can and can't do. **AI Builders gives a team one shared understanding of how AI works, one set of working practices, and the confidence to rebuild its own work around them.**



Teaching DoorDash's analytics organization at its offsite in Seattle, August 2025.

14 cohorts over two and a half years	**3,000+** paid learners	**5.0 / 5** 87 ratings on Maven	**Tencent · Meituan · Xiaohongshu Pinterest · DoorDash** team trainings and internal sessions

More than half of our learners work at leading technology companies, including **OpenAI, Anthropic, Google, Meta, and Amazon.**

FOR THE PEOPLE DECIDING

Rolling out the tools was the easy part.

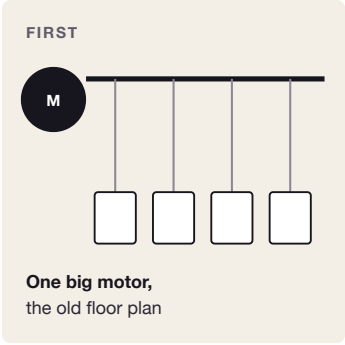
What we hear from the leaders who bring us in, what changes once their team trains together, and where to see the evidence.

	WHAT YOU SEE	WHAT CHANGES	THE EVIDENCE
01	Everyone learned AI on their own, so everyone has a different theory of what it can do.	One shared technical understanding and one set of standard practices: state the outcome, supply the context, set acceptance criteria, check the result. People speak the same language, your few experts stop re-teaching, and work can be handed off without a walkthrough.	**One page of shared rules: 10 → 28 of 28 done right the first time.** Page 5
02	Half the team over-trusts polished output. The other half wrote AI off after one bad try.	A calibrated sense of what AI can and can't do today, reasoned from how it works rather than from headlines, and the habit of telling an answer that looks right from one that is right.	**Our predictions, written down with odds in 2023 and graded since.** Page 6
03	"AI doesn't work on our data."	Most bad results are a method problem, not a model problem: missing context, no definition of done, the wrong layer fixed. People learn to diagnose before they blame, and to put their effort where it pays.	**Same model: 0 of 6 runs right with a 16-step plan, 6 of 6 with five acceptance criteria.** Page 5
04	People are anxious, or quietly opting out.	The fear of no longer being needed, old learning habits, and goals too vague to act on give way to small wins finished early. People learn that crude tools mean the era is early, not that they are bad at this.	**Three versions to get one tool right, and why that is the normal path.** Page 4
05	Everyone is faster. The team isn't.	Workflows, reviews, and metrics priced for an era of expensive execution get recomputed. People own results end to end, and specialists ship tools colleagues can use instead of answering one-off requests.	**Six beliefs from the last era, and why line-by-line approval catches 13.6% of risks.** Pages 3 and 5

How you'll know it worked. A few weeks after training, each person has one finished thing in use, has got unstuck on their own at least once, and can pick up a colleague's project without a walkthrough. Page 8.

Some of your best practices were priced for a different era.

Every best practice carries the price list of its time. Code was expensive, so we reused it. Observation was expensive, so we sampled. Execution was expensive, so we met before acting. AI changed the prices. The hardest part is unlearning: newcomers don't know what works; experts know too well what used to.



When factories first electrified, many swapped the steam engine for one big electric motor and kept the old floor plan: one drive shaft, belts to every machine. Productivity barely moved. The gains came decades later, once factories were redesigned around a small motor at each machine.

Bolting AI onto an old workflow is the big-motor stage. Training has to reach the floor plan.

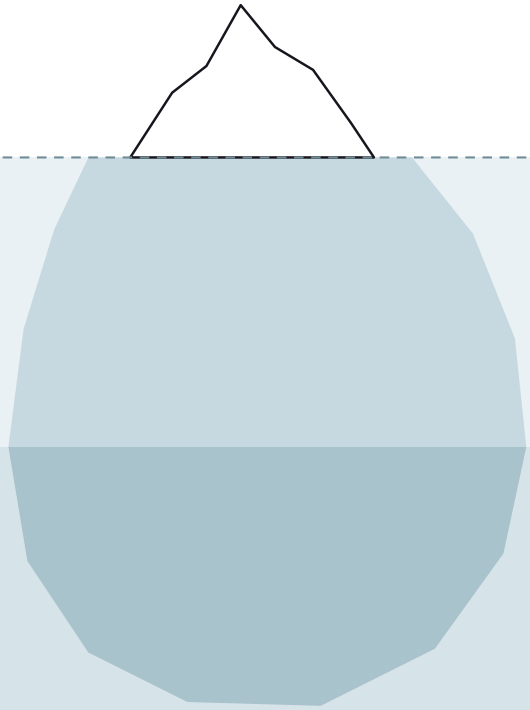
A story from our DoorDash session. The history is documented in Paul David's research on the dynamo and the computer.

	THE OLD BELIEF	WHAT WE SEE NOW
1	~~More AI usage means more value.~~	Activity, tokens, and tool counts all go up. **Dropped into the wrong process, AI amplifies the mess.**
2	~~Faster code is better code.~~	Code still has to serve the business and be maintained. **Generating gets cheaper; getting adopted is what gains value.**
3	~~Individual productivity adds up to team productivity.~~	Everyone feels faster, but **the handoffs, interfaces, and evaluation around them weren't rebuilt.**
4	~~Productivity gains turn into profit.~~	Unit costs fall. **Demand that people will pay for is still scarce.**
5	~~Knowing the concepts means understanding AI.~~	Anyone can explain RAG or MCP. **Nouns are snapshots; real work is continuous,** and that is where the judgment lives.
6	~~Owning a result means checking every step yourself.~~	Owning it now means **checks that scale**: acceptance criteria, independent review, and people on the decisions only people can make.

Beliefs 1–5 come from our session for Tencent's engineering teams; belief 6 from our writing on AI-native work.

Answers are free now. Judgment isn't.

Anyone on your team can ask AI what RAG or MCP means and get a clear answer in seconds. What separates people who get reliable work out of AI sits below the surface, where tutorials don't reach. Training gives a team the same picture of that layer, and the practice to build it.



ABOVE THE SURFACE · FREE AND COPYABLE

Prompts, tools, tutorials, RAG, MCP, agents

AI explains them in seconds. Anyone can copy them. They change every month.

WATERLINE

BELOW · WHAT DECIDES THE RESULT

Technical intuition

What to give AI, where it will break, what to check.

Product intuition

What is worth building, for whom, and what good looks like.

DEEPEST · WHERE THAT JUDGMENT COMES FROM

A world model

Your grasp of causes, incentives, and constraints, sharpened by predictions that reality can grade.

Craft

What you can see before you can explain it. Built only by doing.

"Knowing what RAG is lives above the water. Knowing you don't need it this time lives below."

Why it has to be built, not watched

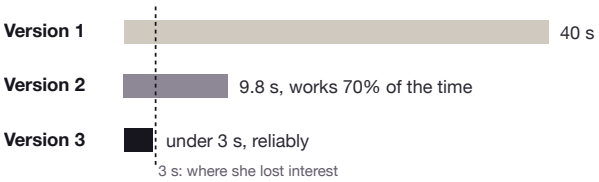
Above the waterline, every question has a standard answer, which is exactly what AI is best at. Below it, there is no standard answer to hand over. Using AI well is a craft, like swimming: you can sit beside Michael Phelps through ten thousand hours of training and still sink when you get in the water.

So every learner builds on a real problem, with someone beside them who has seen the failure before.

What stops people, and what gets them moving

Three things hold people back: old learning habits, goals too vague to act on, and the fear of no longer being needed. What helps is concrete: a small loop finished this week, shown to someone, discussed. And knowing that first versions are meant to be thrown away.

A learner's reading assistant for his five-year-old: seconds per reply



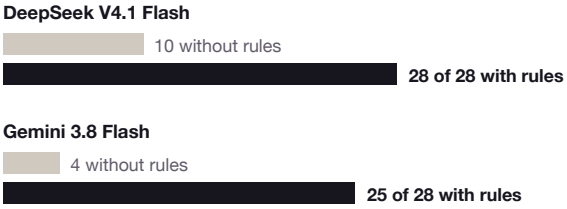
CLAIMS WE TEST BEFORE WE TEACH

We show our numbers.

Our course text is built on controlled comparisons: change one thing, hold everything else, grade against an answer key. Four results that matter to a team.

One page of shared rules

28 voice memos, each tidied in a fresh window with the same one-line request. The only change: a five-line rules file. Right first time only if all four checks pass.

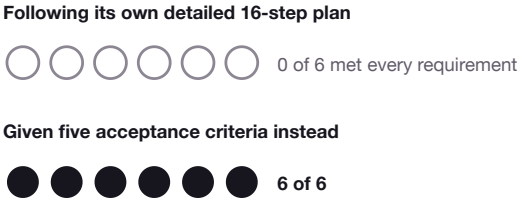


Standard practices aren't bureaucracy. Written down once, they make every run more reliable, on more than one model.

AI Builders course text, lesson 2.4, November 2026.

Describe the destination, not the steps

A 55-minute meeting recorded in two files that overlap by 40 seconds, with a quiet phone participant. Same top model, six runs each.

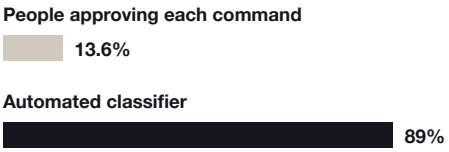


Every check in a step list checks the list's own assumptions. All six plans assumed the files didn't overlap. Given only the destination, the model found the overlap itself.

AI Builders course text, lesson 3.7.

Step-by-step approval is false control

Anthropic blind test, 1,053 participants: dangerous commands slipped into an AI agent's session. Share caught:



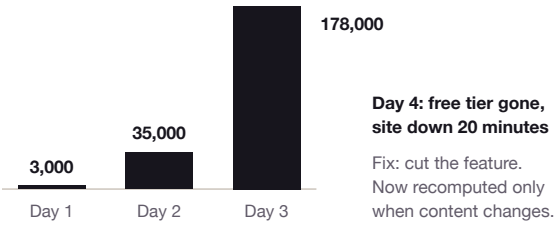
People caught about 17% early in a session, and 5% after 50 approval prompts.

Put people on direction and plans, not on clicking approve. Shown one command at a time, people approve 97%; shown a whole plan, they reject 39%.

Anthropic data, cited in AI Builders lesson 3.8.

The best fix was a subtraction

Ask Lizheng, Yuzheng's production Q&A site, after it launched a personalized "not yet seen" list. Database requests per day:



Its cost was questions × visitors. Once you see that, buying more capacity isn't the fix. Seeing it is the skill.

AI Builders Cohort 14, session 1.

Hype we checked. When multi-agent setups became the default advice, Yan worked through the public experiments: on one benchmark a team of agents added 7.3 points, while simply raising one model's reasoning effort added 11.4. His lecture on it is free to watch.

JUDGMENT ON THE RECORD

We write our calls down, with odds.

In March 2023, before GPT-4 launched, Yuzheng published six bets about the next two years, each with a probability. Here is how reality graded them, plus one call from our January 2025 course.

	THE BET	ODDS	RESULT	WHAT HAPPENED
1	Expert-level answers to factual questions	99%	Yes	o1 beat PhD experts on a science test (2024)
2	Expert-level answers to judgment questions	80%	Mostly	Half of messages seek information or advice
3	A digital copilot for everyone	70%	Yes	800 million weekly ChatGPT users (2025)
4	10× programming; AAA games in six months	60%	Partly	Coding agents, yes; AAA in six months, no
5	An AI apprentice that out-earns Yuzheng	50%	Not yet	
6	An article full of original knowledge	5%	No	His book still needed him for that
’25	Agents are next: an 18-month window (our Agentic AI course, January 2025)		Yes	Claude Code and Codex became the main way builders work

The likely bets landed; the long shots didn’t. What we got wrong, in Yuzheng’s own 2025 note: he overestimated OpenAI, and how fast models would improve. Original: yuzheng.substack.com

2023 · YUZHENG

“Inference cost will drop by more than two orders of magnitude.”

GPT-4 cost $30 per million tokens in 2023. GPT-4o mini, which beat it in chat, cost $0.15 in 2024.

2023 · YUZHENG

“Its best way to use knowledge is search, not adding knowledge to the model.”

Retrieval and context became the default. In the course: models are rented, context is owned.

MARCH 2025 · YAN

“The first step in learning agentic AI is to forget every framework.”

Sixteen days later he published “The underrated Claude Code.” Claude Code later took off.

Two kinds of judgment, taught together

What is true → what is worth doing → how to make it happen reliably



Yuzheng Sun, PhD

What is worth building, and who owns the result.

Cornell PhD in Economics. Economist at Amazon, data scientist at Meta, Vice Director of Data & AI at Tencent IEG, Principal Data Scientist at Statsig (later acquired by OpenAI). Co-organizer of AI Engineer Shanghai.



Yan Wang, PhD

Why models work, why they fail, and how to make systems reliable.

Columbia PhD in Electrical Engineering. Nearly 40 papers across CVPR, NeurIPS, KDD, and other venues; 3,000+ citations. Creator of the open-source devin.cursorrules.

Fourteen cohorts over two and a half years · live sessions rebuilt for Cohort 14, October 2026 · course text rewritten in English and Chinese, November 2026 · lifetime updates included.

SEE IT BEFORE YOU DECIDE

Watch the sessions we ran inside other companies.

Recordings of our team trainings and internal sessions are free to watch. Pick the one closest to your team and judge the teaching yourself.

DoorDash Analytics offsite	How to manage context so AI does high-quality work Why AI fails on undocumented data, and how analysts become navigators instead of rowers.	25 min · English
Pinterest Data Science Weekly	AI breaks data science's impossible triangle Speed, scale, and explainability used to be pick-two. A live build on synthetic data shows AI delivering all three.	58 min · English
Tencent Engineering series	Common mistakes in an engineering team's AI transition Why individual speed doesn't become team speed, and what has to be rebuilt.	1 h 57 min · Chinese
Meituan Top Talk	From AI user to AI builder Mindset over habits, habits over methods, methods over tools, and why the order matters.	1 h 55 min · Chinese
Xiaohongshu Non-technical teams	From 30% faster to 10× For operations, content, commercial, product, and brand roles, built from the company's own survey of 16 problems.	1 h 34 min · Chinese

Also public: talks at Columbia University and the Yale Center Beijing; a conversation with The Pragmatic Engineer; "Data Analytics in the Age of AI" with speakers from OpenAI, Anthropic, Gen, Lyft, and DoorDash. Full list: superlinear.academy/c/open-videos

"There are many courses teaching you specific tricks but this course helps you build the right mindset and empower you to teach yourself more effectively."

Shuyang · Member of Technical Staff, **OpenAI** · former Sr. AI Manager, Uber

"In the era of AI, everyone becomes a manager… Accessing AI tools does not guarantee results. Just like not every manager is a good one, not every AI-adopter is an effective user."

Wang · Machine Learning Engineer, **GEICO**

"I waited a couple of months before sharing feedback because I wanted to see how things actually held up in practice… I've been able to apply what I learned here directly to my job."

Ying · Principal Engineer

"There can be different new tools coming, but principles remain!"

EZ · Engineer, **Anthropic**

Public Maven reviews, shortened with ellipses. Employers identify the reviewers and do not imply company endorsement.

WAYS TO WORK TOGETHER

Start with the simplest step that solves the problem.

01	**Working session**	One real question with your leads: what result you want, which assumptions to test, and what the right next step is.	**From $2,000**
02	**Team seats in a scheduled cohort**	Your people join a live Maven cohort alongside professionals from other companies. Next cohort: Feb 15 – Mar 1, 2027.	**$1,999 per learner** 2–9 learners: 20% off ($1,599.20) 10+ learners: 25% off ($1,499.25)
03	**Private cohort**	The same course for your organization only, on dates we set together.	**From $44,977.50** 30+ learners · $1,499.25 each
04	**Course customization**	Your cases, data, and exercises, when they change what the team needs to learn.	**From $20,000** in addition to tuition
05	**Fully custom program**	We study the roles, workflows, materials, and evaluation first, then design the program, the delivery, and the tools your team keeps.	**From $100,000**
+	**Self-paced plan**	AI Builders and AI Architect in English, plus all four courses in Chinese. No live sessions. Every future update is included, so buying now includes the Nov–Dec 2026 refresh. With 12 months of Stay Superlinear membership.	**$1,048 per learner** $1,299 from Jan 1, 2027

COHORTS INCLUDE Two 2-hour live workshops and an optional 90-minute office hour · five guided projects · a capstone of 5–30 hours with feedback · lifetime access to 100+ lessons and all future updates · certificate of completion · one year of Stay Superlinear membership.

The test we hold ourselves to

1. Does each person have one finished thing that is in use?
2. Has each person, at least once, got unstuck on their own?
3. Can someone pick up a colleague's project and change it without a walkthrough?

Ask these a few weeks after training. If the answers are no, the training didn't do its job.

Tell us what should be different afterwards.

Send the team, the problem today, the result you want to change, and rough timing. We'll recommend a starting point and send a quote.

Email yz@superlinear.academy
Course maven.com/superlinear/aibuilders
Formats lizheng.ai/en/collab/enterprise

Prices in USD, before applicable taxes. Cohort prices, team discounts, and dates follow Maven's settings as of November 12, 2026. The self-paced price changes at 00:00 on January 1, 2027, Beijing time; final terms are shown at checkout, and invoices for approval are available on request. Private cohorts, customization, and custom programs are subject to instructor availability and a mutually executed order form that sets final scope, schedule, payment, cancellation terms, and taxes. Quotes are valid for 30 days. Companies named are public team trainings and internal sessions; custom engagements are named only with client permission. Experiments cited from our course text use the models named at the time of writing.